

CSE 676: Assignment #1

Due data: 11:59 PM, March 18th, 2023

0.1 Softmax [20 points]

- 1) [10 point] Prove that softmax is invariant to constant shifts in the input, *i.e.*, for any input vector \mathbf{x} and a constant scalar c , the following holds:

$$\text{softmax}(\mathbf{x}) = \text{softmax}(\mathbf{x} + c) ,$$

where $\text{softmax}(\mathbf{x})_i \triangleq \frac{e^{x_i}}{\sum_{i'} e^{x_{i'}}$, and $\mathbf{x} + c$ means adding c to every dimension of \mathbf{x} .

- 2) [10 point] Let $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{c}$, where \mathbf{W} and \mathbf{c} are some matrix and vector, respectively. Let

$$J = \sum_i \log \text{softmax}(\mathbf{z})_i .$$

Calculate the derivatives of J w.r.t. \mathbf{W} and \mathbf{c} , respectively, *i.e.*, calculate $\frac{\partial J}{\partial \mathbf{W}}$ and $\frac{\partial J}{\partial \mathbf{c}}$.

0.2 Logistic Regression with Regularization [20 points]

- 1) [10 point] Let the data be $(\mathbf{x}_i, y_i)_{i=1}^N$, where $\mathbf{x}_i \in \mathbb{R}^d$ and $y_i \in \{0, 1\}$. Logistic regression is a binary classification model, with the probability of y_i being 1 as:

$$p(y_i = 1; \mathbf{x}_i, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}^T \mathbf{x}_i) \triangleq \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}_i}} ,$$

where $\boldsymbol{\theta}$ is the model parameter. Assume we impose an L_2 regularization term on the parameter, defined as:

$$\mathcal{R}(\boldsymbol{\theta}) = \frac{\lambda}{2} \boldsymbol{\theta}^T \boldsymbol{\theta}$$

with a positive constant λ . Write out the final objective function for this logistic regression with regularization model.

- 2) [10 point] If we use gradient descent to solve the model parameter. Derive the updating rule for $\boldsymbol{\theta}$. Your answer should contain the derivation, not just the final answer.

0.3 Derivative of the Softmax Function [30 points]

- 1) [10 point] Define the loss function as

$$J(\mathbf{z}) = - \sum_{k=1}^K y_k \log \tilde{y}_k ,$$

where $\tilde{y}_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}$, and (y_1, \dots, y_K) is a known probability vector. Derive the $\frac{\partial J(\mathbf{z})}{\partial \mathbf{z}}$.

Note $\mathbf{z} = (z_1, \dots, z_K)$ is a vector so $\frac{\partial J(\mathbf{z})}{\partial \mathbf{z}}$ is in the form of a vector. Your answer should contain the derivation, not just the final answer.

- 2 [10 point] Assume the above softmax is the output layer of an FNN. Briefly explain how the derivative is used in the backpropagation algorithm.
- 3) [10 points] Let $\mathbf{z} = \sigma(\mathbf{W}^T \mathbf{h} + \mathbf{b})$, where $\sigma(\cdot)$ is the sigmoid function, \mathbf{W} is a matrix, \mathbf{b} and \mathbf{h} are vectors. Use the chain rule to calculate the gradient of \mathbf{W} and \mathbf{b} , *i.e.*, $\frac{\partial J}{\partial \mathbf{W}}$ and $\frac{\partial J}{\partial \mathbf{b}}$, respectively (it is enough to derive the gradients for one element of the matrix/vector parameter).

0.4 MNIST with FNN [30 points]

- 1) [10 points] Design an FNN for MNIST classification. Draw the computational graph of your model.
- 2) [20 points] Implement the model and plot two curves in one figure: *i*) training loss vs. training iterations; *ii*) test loss vs. training iterations.
 - You can use online code. However, you must reference (cite) the code in your answer.
 - Submission includes the plot of the two curves and a link to the runnable code in your Github account (in the Github, you need to write a ReadMe file containing instructions on how to run the code).
 - You can use any packages, but Pytorch is recommended.